

Same company. Same day. Two different answers.

Its published ESG grade puts it in the bottom fifth of its own basket. Fourteen dated filings, exchange notices and news items say it is the fastest-improving name in it.

+0.80

DISAGREEMENT

ITS RATING SAYS

18th of 52

grade BB · bottom fifth

THE EVIDENCE SAYS

98th percentile

14 dated sources · 2023–2026

A rating tells you where a company **stood**.

Not which way it is moving, and not what would change the answer. ESG momentum is a real edge — here is the half its own author reports too.

A level, published late

A score in a file closed months ago, refreshed on the rater's cycle — not the company's.

No direction

Nothing in the number tells you whether the company is improving or deteriorating right now.

Not challengeable

You cannot ask a rating which piece of evidence is holding it up, or what would change it.

+55.1%

the basket, cumulative — against +6.4% for MSCI AC ASEAN

-4.6%

the same basket, year-to-date 2025, on a ~45% banking weight

28.9%

probability an ESG improver beats the index at one year

We rank all 52 companies **twice**.

Once on the rating the market already uses. Once on dated public evidence. Then we report the gap — with a direction, a date on every source, and a list of what would change it.

Thinks

It interrogates before it shows: four short questions that set the filters, the risk tier and how fast old news stops counting. Every answer changes what is on screen.

Challenges

A loose framing comes back sharper. It pushes back on the question instead of answering the wrong one.

Competes THE DIFFERENTIATOR

It takes a position against a named rating, on evidence the rating cannot see — and shows its working.

24

both agree — and we say so

21

improving, already rated for it

4

rated behind the evidence

2

flattered by an old score

1

poorly rated, still slipping

Ninety seconds. Six clicks. Same answer.

0:00	The board	52 real companies, one deterministic run — the id is top-right. Same files in, same answer out.
0:15	Disagreement matrix	Left–right: what the rating thinks. Bottom–top: what the evidence says. Hollow dots are companies with no evidence — we draw the absence.
0:30	The RHB chip	Rated 18th. Evidence 98th. Fourteen dated sources. One click, no hunting for the dot.
0:45	Evidence trail	Every signal dated, sourced, directional, with the reason it was weighted that way. Three clicks to the primary source.
1:05	What would change this	Pull any one source out and we re-score, and show which removals flip the verdict.
1:20	Verify this evidence	Recompute every hash against the root anchored on a public chain. MATCH — proof on-chain, no licensed data off the machine.

No chain endpoint? The run says anchor_pending. It never claims to be anchored when it is not, and the check still runs locally against the stored record.

We wrote the bar down **before** the evidence.

A parameter may change for a reason that can be argued without reference to the answer it produces. It may not change because we dislike the answer.

PRE-REGISTERED · 25 AUG 2026 · BEFORE THE SWEEP

Deterministic

No clock, no randomness, no model in the scoring path. Same files, same run id, same quadrants.

Five guards

No date or no working source, the fact is dropped. We decide the source type from the domain — never the publisher.

The cap that costs us

Company-published material is capped at 0.5 confidence. 86% of a public search is company-published.

0

thresholds moved to produce that list

4

names crossed the unmoved bar, on better sources

25

signals on the best-covered company — the bar is reachable

1

published failure case kept in the backtest

The AI reads. The rules score.

	THEY DO	WE DO
A rating	Publishes a level, late	Publishes a gap, with a direction and a date on every source
A stock picker	One blended number	Two axes, never merged — ESG evidence, and a separate money gate
Chain ESG platforms	Issuer-side: the company writes its own data to a chain	Investor-side and adversarial — the company's own word is our least-trusted source
“AI-scored ESG”	A model assigns the score	The model only gathers; scoring is rule-based and costs zero tokens

UOB clears the ESG bar and fails the money read — earnings down 22.6%. We show both and never average them. One combined score would have buried it.

S\$0.079 per company per year. Measured.

S\$4.12

a year to cover the whole 52-company basket

S\$39.65

a year at 500 companies

S\$158.60

a year at 2,000 companies

S\$118.45

total spend to build what you just saw

Zero marginal scoring

Rules route the evidence and deterministic code scores it. No language model is reachable from either — per-signal token cost is structurally zero.

All AI cost is gathering

Charged per sweep of a company, not per signal. Running sweeps off-peak is an exact 50% saving.

People, not tokens

Metadata verification is hand-done. Day rates come from live Singapore job postings — ESG analyst S\$290/day — and are flagged as a floor.

Pilot: S\$34,965 people cost

The pilot total is stamped INCOMPLETE, because the data licence and compliance lines are still blank.

No falling cost-per-company curve until a licence quote exists. Inference is linear in companies, so on this cost alone that line is flat. Two inputs are on the slide as blanks: a baseline-data licence quote, and one timed batch of ten companies.

What it does **not** do.

On the slide, before anyone has to ask.

86% is self-published

Most of what a public search returns is written by the companies themselves, and we cap how far that can ever count. That is why the top quadrant holds four names, not forty.

Documentary, not telemetry

We read filings, regulator actions, exchange notices and press. We do not meter a smokestack.

2 of 52 have no evidence

Drawn hollow on the matrix and counted. An absence of evidence is a finding, not a hole to fill.

11 of 52 are financially weak

And 2 are unreadable — reported as unknown, never marked down for our missing data.

Licensed raters stay off-screen

MSCI, S&P and Sustainalytics cannot be shown — licence, not capability. LSEG is fetched live because it is the one free keyless endpoint.

Two delisted names stay in

Badged, and excluded from investable output by rule. A static list going stale is the argument.

It never says buy or sell. It says where it **disagrees** — and exactly what would change its mind.

52

ASEAN companies, one verified research basket

391

dated, sourced events behind the verdicts — 48 of 52 covered

\$0.079

AI cost per company per year, measured

0

thresholds moved to produce our list

The bar was fixed before the evidence existed, and the quadrant filled because we found better sources — not a smaller number.